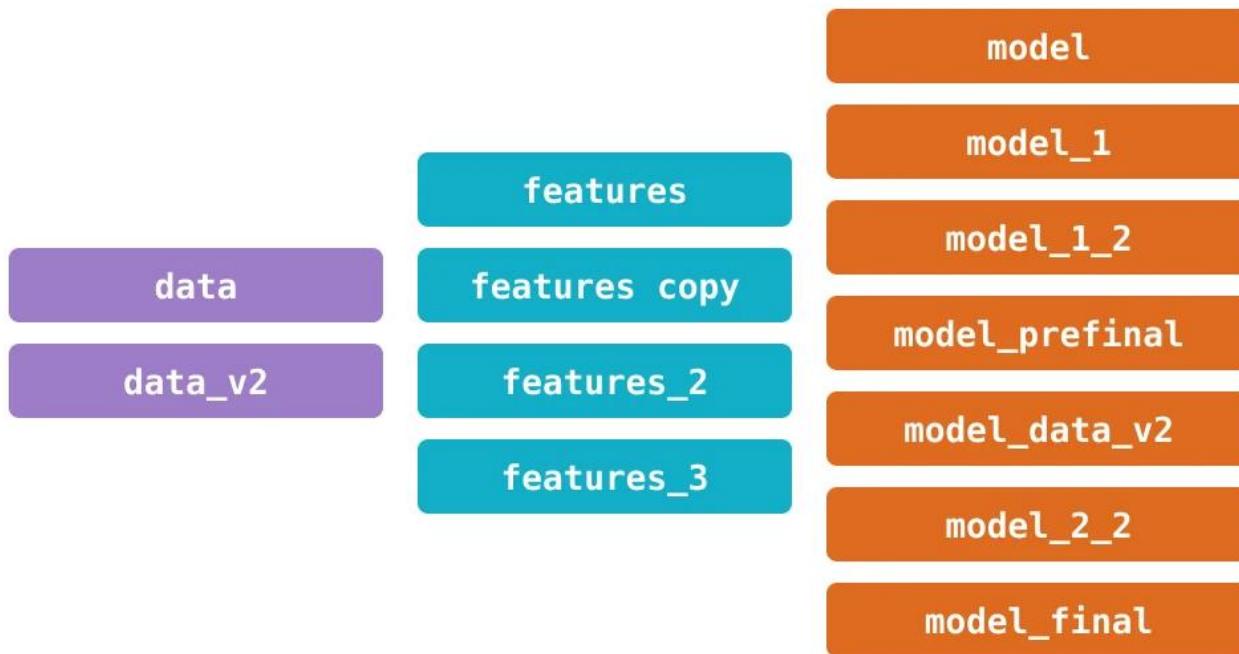


Versioning

Florian Bär
MLOps

Information Technology
09.10.2025

Complexity of data science projects



Overview

- Version control for
 - Code
 - Data
 - Models
- Collaboration on
 - Code
 - Data
 - Models

Version control for code



mercurial



git

Version control for code

- Go-to versioning system: Git



Version control for code

- Go-to versioning system: Git
- [HSLU Git course on Ilias](#) (self-study)
- [Git introduction on GitHub](#) (explains the most important commands)
- Go-to code hosting platform: [GitHub](#)
 - Open-source alternative for self-hosting: [GitLab](#)
- [Introduction course to GitHub](#)
- [Get started tutorial](#)

Version control for code

- Different versions are saved in *commits*
- Difference from one commit to the next is the *diff*
- Code is just text, so storage remains rather concise
 - Don't store data/models/logs/plots in vanilla Git
- See the commit history with `git log`, or the last commit including the diff with `git show`
- Each commit is identified by a *commit ID*
 - Can print the commit ID of the last ("HEAD") commit with `git rev-parse HEAD`

Version control for code

- Important to reproduce experiments: Know which version of the code you used
- Simple solution: Print the current commit ID to a file in your experiment directory:

```
git rev-parse HEAD > PATH_TO_EXPERIMENT/commit.log
```

- Experiment tracking tools can help you with tracking code versions, too
 - See e.g. [Weights and Biases](#)

Version control for code

- You do not want to track all files in Git together with the code
 - Data
 - Logs
 - Result files
 - Model checkpoints
 - OS/IDE files
- Add those files to the `.gitignore` file, either individually or with a regular expression pattern:

`*.csv` (all csv files)

`logs/` (the directory "logs" and all its files)

`data_*/` (all directories starting with "data_")

Version control for code

- Git Hooks
- Tools (lefthook, husky)
- post-checkout / post-update / **pre-commit** / prepare-commit-msg / pre-push / pre-rebase
 - E.g. Lint your code, prevent push .env, run tests (not recommended)
- <https://github.com/CompSciLauren/awesome-git-hooks>
- [prevent-bad-push](#) (prevents you and others to commit “WIP” messages)

Version control for data

- Exact dataset of an experiment needs to be tracked for reproducibility
- Version control is different from code
 - Larger file sizes
 - Tracking and storage are separated
 - New data points come in more frequently than code changes, but dataset versions typically change less frequently
 - Diffs have different semantics
- Tools
 - [Git LFS](#)
 - [DVC](#)

Version control for data: Git LFS

- Determine which file types to track:

```
git lfs track "*.csv"
```

(or add a line with *.csv to the .gitattributes file)

- Use git to track changes (add/edit/remove)

```
git add .gitattributes *.csv
```

```
git commit -m "Tracking all .csv files with Git LFS"
```

- Data itself is stored separately from code

- GitHub has an integrated solution (free version for up to 500MB)
- [List of alternatives](#)
- Run your own by adapting the [reference implementation](#)

Version control for data: DVC

- Alternative to Git LFS
- Also uses Git for version control
 - Stores file hashes in ".dvc" files and tracks those with Git
 - Stores files themselves in cache directory
- Can attach to any storage (on-premise or cloud)
 - More [supported storage options](#) than Git LFS
- Can be used to track models, too

Version control for data: DVC

- Assume: data in "data" folder, model checkpoint in "checkpoint.pt"
- Add data and model to version control:

```
dvc add data checkpoint.pt
```

- Moves the files into the cache (.dvc/cache) and adds them to .gitignore
- Creates files data.dvc and checkpoint.pt.dvc that contain file hashes that point to the cached files
- data.dvc and checkpoint.pt.dvc are tracked in Git:

```
git add data.dvc checkpoint.pt.dvc
git commit -m "First model"
git tag -a "v1.0" -m "model v1.0"
```

Version control for data: DVC

- Assume: new data added to "data" folder, retrained model → new checkpoint
- Add the new data and model to version control:

```
dvc add data checkpoint.pt  
git add data.dvc checkpoint.pt.dvc  
git commit -m "Second model, trained with more data"  
git tag -a "v2.0" -m "model v2.0"
```

- Can easily switch back to old state:

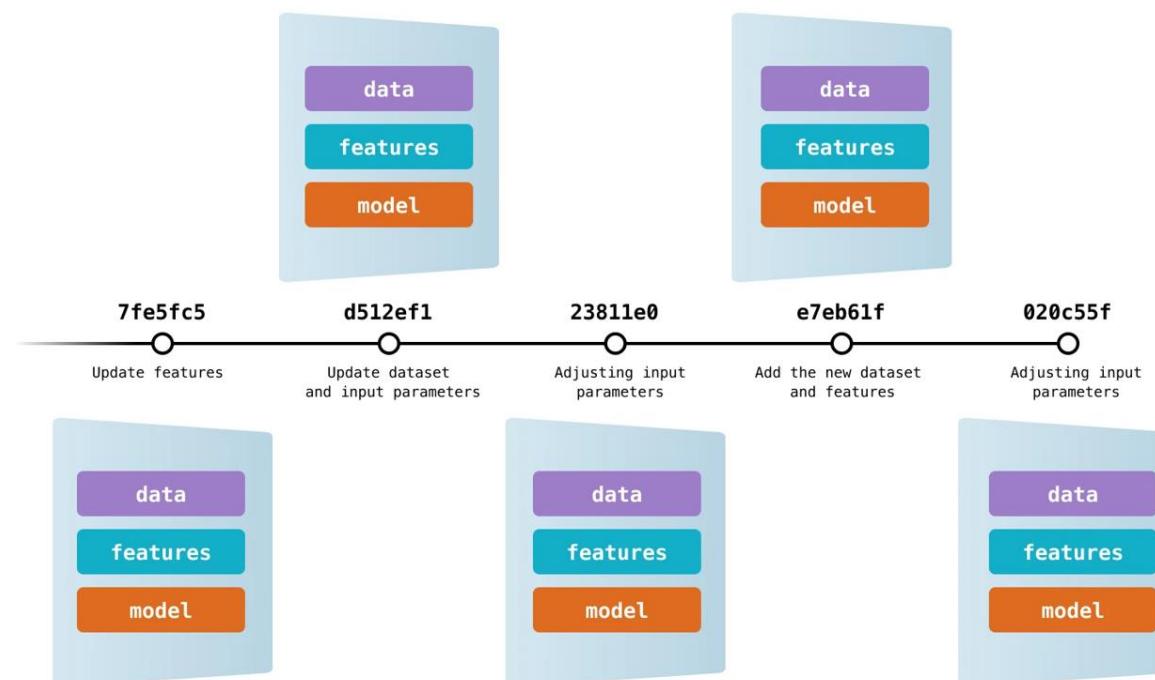
```
git checkout v1.0  
dvc checkout
```

- Or just switch to old data and keep rest:

```
git checkout v1.0 data.dvc  
dvc checkout
```

DVC: Timeline of experiments

- Single timeline for versions of data and model



Version control for models

- Important for reproducibility
- Manual
 - Track on running an experiment
 - Save code/data version, all hyperparameters, exact command used
 - Easier to make mistakes or become inconsistent over time
 - Harder to follow while collaborating in a team
- Automatic
 - Weights and Biases: [model registry](#)
 - MLFlow: model registry
 - DVC: Same as code, as seen before

Version control for pipelines with DVC

- Create a stage and adds it to the dvc.yaml file:

```
dvc stage add \
-n train \  (name)
-d train.py -d data \  (dependencies/inputs)
-o checkpoint.pt -o training_log.txt \  (outputs)
-m metrics.csv \  (metrics -> key/value pairs)
python train.py --model cnn    (command)
```

- The stage description (dvc.yaml) and the inputs/outputs (via DVC) can be added to Git version control (git add & git commit)
- Reproduce the outputs of a stage:

```
dvc repro train
```

Collaboration

Collaborating in data science projects

- MUST agree on a shared process
 - Project structure
 - Storing and sharing of data and models
 - Experiment tracking
 - Naming standards, Code formatting
- Using tools can serve as standardization mechanisms
 - If a tool enforces a process, all team members will follow the same process
 - Examples:
 - Tools mentioned earlier

Collaborating on code

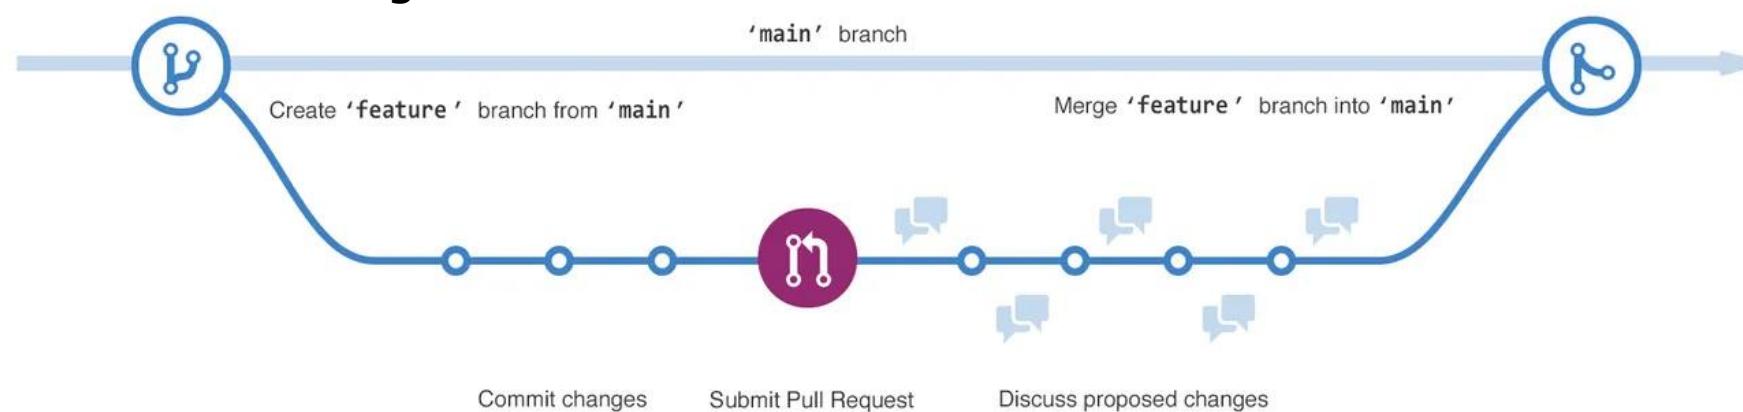
Organize code with branches:

- Typically: `main` (previously `master`) branch holds the current (working) state of the project
- Different setups are possible
 - "development" branch for dev, "production" branch for live code
 - Even possible to have separate repositories for different projects
 - Makes it harder to share code (e.g. utils)
- Tool: `git-flow`
 - Feature, Releases, Hotfixes, Tags

Collaborating on code

Features are implemented in their own branches

- New branch is created from a specific commit of the base branch
- Changes are applied
- When a feature is ready to be merged back (testing!), open a pull request
- Collaborators discuss if further changes are necessary (*code review*)
 - Can lead to new changes being committed
- Feature branch is merged back into base branch



Collaborating on code

- In the pull request, all the changes (diff) to the base branch are shown
 - Red: removed, green: added



The screenshot shows a GitHub pull request diff interface for a file named README.md. At the top, it says "Showing 1 changed file with 3 additions and 3 deletions." There are "Split" and "Unified" view options. The diff itself shows the following changes:

```
@@ -1,3 +1,3 @@
- # test-area-2
- edit1
- edit2
+ # About me
+
+ My name is Mona Lisa.
```

The first three lines are marked as deleted (red background), and the last three lines are marked as added (green background).

Collaborating on code

- Standard practice: Code reviews
 - Frequency depends on the company:
For all code, or just for certain features/when requested
- Code reviews are not easy
 - Expectation management
 - Refactorings? Code style? Unit tests?
TONE?

> RISC-V Patches for the 6.17 Merge Window, Part 1

No. This is garbage and it came in too late. I asked for early pull requests because I'm traveling, and if you can't follow that rule, at least make the pull requests *good*.

This adds various garbage that isn't RISC-V specific to generic header files.

And by "garbage" I really mean it. This is stuff that nobody should ever send me, never mind late in a merge window.

Like this crazy and pointless `make_u32_from_two_u16()` "helper".

That thing makes the world actively a worse place to live. It's useless garbage that makes any user incomprehensible, and actively *WORSE* than not using that stupid "helper".

Collaborating on code

- Common Terms
- Definition of Ready (DoR)
 - Work is ready to enter the Sprint.
- Acceptance Criteria (AC)
 - Team works to meet the **unique user functionality** required.
- Definition of Done (DoD)
 - Work meets the **universal quality standard** and is **Done** (releasable).

Collaborating on code

- Best practices
 - Minimal edits: Only change what is absolutely requested and part of the current change request
 - No need to rename local variables, reformat, or do refactoring → Do these tasks in a separate change
 - Style/formatting guides: Using a shared formatter avoids many changes due to different formatting in different code editors
 - Don't just look at superficial changes ("Do I like this variable name?") but at functionality ("Can this code result in a deadlock?")
 - Be nice when writing a review! It can be hard to have our work criticized, so do it in a constructive and positive manner.
 - Don't be overconfident. The author of the code has usually spent a lot longer thinking about it than you have.

Collaborating on code

Concurrently working on the same code:

- Try to separate the changes into smaller parts
 - Use feature-branches
- Git is good at resolving conflicts, but not perfect
 - If the edits overlap, you will have to resolve conflicts manually
 - If not, there is still a chance that a formatting change screws up automatic merging of the edits

Collaborating on code

- *Releases* are snapshots of the current state of the codebase
 - Stable versions of the product
- Makes it easier to distribute code to others
 - Can summarize changes in the release notes
 - Code is packaged with release notes and binary files
- Semantic Versioning (guide to naming versions)
- Soft version of a release: *tag*
 - Marks a specific commit ID with a tag
 - Description explains the reason for the tag

The screenshot shows a GitHub repository interface. At the top, there are two commits listed:

- emnlp2022** · a70e9cdc · Added bart encoder attention base figure. · 1 year ago · EMNLP 2022 submission: Hallucination detection
- emnlp2021** · ca6079f2 · Small layout changes. · 1 year ago · EMNLP 2021 workshop camera-ready: Sentence planner

A blue arrow points from the text "My tags for my paper repository" at the bottom left towards the "emnlp2022" tag.

Collaborating on data

- Use a common process, or a common tool that comes with a standard process
- Determine preprocessing procedures
 - Run preprocessing once and store the results
 - Do not run preprocessing for every experiment!
- Define results file format
- Set up automated pipelines
- Automatically analyze the results
 - E.g. generate a graph from a metrics.csv file
 - The less manual effort, the more time for understanding the problem

Collaborating on models

- Create model releases
- Central model registry
 - Weights and Biases
 - DVC
 - Hugging Face
 - MLFlow
- Integrating model versioning with experiment tracking helps
- Locally run experiments may as well not exist

AI Jobs Report 2025



HSLU Lucerne University
of Applied Sciences
and Arts

x28

LAC

The Applied AI Center

AI JOBS REPORT 2025

An Overview of the
AI Jobs Market in Switzerland



- An overview of the Swiss AI Jobs Market