

# Containerization

Florian Bär  
MLOps

**Information Technology**  
16.10.2024

# Why?

- Any ideas?



works on  
my  
machine



# Portability, Consistency, Scalability

Different machines have different

- Hardware
- Operating system
- Installed software

Goal: Application should run the same way on different machines

- Locally (for development)
- On-premise/private cloud
- Public cloud (different providers)

# Overview

- Dependency management
- Containerization
  - Docker, Podman, ...
- Managing containers
  - Kubernetes, Openshift, ...
- Building pipelines
  - Airflow, Kubeflow, Metaflow, TFX

# Dependency Management

- Libraries help us develop code faster and safer
  - Reuse existing functionality
  - Rely on thoroughly tested and maintained code
- But: Introduces additional complexity
  - Interfaces, Adapters etc. needed.
  - Everyone needs to use the same library versions in a project
  - in Python?!
- Best practice: Set up an environment (a set of all libraries and their versions) for each project

# Python Dependency Management

- Install Python libraries from
  - [Pip](#) (package installer for Python)
  - [Conda](#) (multiple languages)
  - Source

PyTorch Build	Stable (2.1.0)	Preview (Nightly)		
Your OS	Linux	Mac	Windows	
Package	Conda	Pip	LibTorch	Source
Language	Python	C++ / Java		
Compute Platform	CUDA 11.8	CUDA 12.1	ROCM 5.6	CPU
Run this Command:	<pre>conda install pytorch torchvision torchaudio pytorch-cuda=11.8 -c pytorch -c nvidia</pre>			

# Python Dependency Management

- Install Python libraries from
  - [Pip](#) (package installer for Python)
  - [Conda](#)
  - Source
- Environments
  - Machine environment (you don't want to go down this road)
    - Needs root access
  - User environment (neither this one)
    - `pip install --user torch`
  - Virtual environment (does the job, but not always – Win, Mac, Lin specific)
    - `virtualenv`/`venv`, `pipenv`
    - Conda (Anaconda/Miniconda)/mamba
    - Poetry
    - uv

# Sharing Dependencies

- pip
  - One configuration file per project specifies the environment:
  - `requirements.txt`
    - Contains library name and version (`torch==2.0.1`)
    - `pip freeze > requirements.txt` (save)
    - `pip install -r requirements.txt` (recreate)
    - But: only includes libraries installed with pip
- Poetry, uv
  - Same, but with lock file
  - Fixed versions with hashes

# Sharing Dependencies

- Conda
  - environment.yml
    - Export environment from conda: `conda env export -f environment.yml`
    - Can recreate the same environment on a different machine:  
`conda create -f environment.yml`
    - Includes channels from where libraries were downloaded, and dependencies (pip and non-pip)
    - [How to manage environments with conda](#)


# Conda Demo

# Containerization

# Docker



- Isolates the application from the hardware/OS
- Command-line interface (CLI) based
  - Good for automation/scripts
  - GUI: Docker Desktop
- More lightweight than virtual machines
  - Does not include the OS



# Docker

- Container
  - Sandboxed process, isolated from all other processes on machine
- Image
  - Package of software used by the container
  - Must contain all dependencies, scripts, data
- [Docker Hub](#): Publicly available Docker images
- [Docker playground](#)

# Dockerfile

- The Dockerfile specifies all operations performed when starting the image
- Example:

```
# syntax=docker/dockerfile:1

FROM node:18-alpine
WORKDIR /app
COPY . .
RUN yarn install --production
CMD ["node", "src/index.js"]
EXPOSE 3000
```

where to look for  
the Dockerfile

- [Dockerfile reference](#)
- Build the image with: docker build -t tag .

# Running a Container

- Run an image with:

```
docker run -d -p 127.0.0.1:3000:3000 tag
```

- -d: detach (run in background)
- -p: publish (port mapping between host and container)
  - Format is *host:container*
  - Necessary to access the application from the host

- To list all the containers: docker ps (-a)

# Further Useful Commands

- `docker stop container-id` (stops running container)
- `docker rm container-id` (removes stopped container)
- `docker tag old-name new-name` (rename image)
- `docker push new-name` (push image to Docker Hub)
- `docker exec container-id command` (execute a command in a running container)
- `docker volume create volume-name` (mount a volume for persistent storage)
- `docker logs -f container-id` (see log of a container)

# Communicating Between Containers

- Containers are isolated processes
- Can allow them to communicate by placing them on the same network: `docker network create network-name`
- Connect containers to network at startup or during runtime



# Docker Compose

- Define and share multi-container applications
  - Defined in a YAML file

- Start the full application with docker compose up and stop it with docker compose down

compose.yaml

```
services:  
  app:  
    image: node:18-alpine  
    command: sh -c "yarn install && yarn run dev"  
    ports:  
      - 127.0.0.1:3000:3000  
    working_dir: /app  
    volumes:  
      - ./:/app  
    environment:  
      MYSQL_HOST: mysql  
      MYSQL_USER: root  
      MYSQL_PASSWORD: secret  
      MYSQL_DB: todos  
  
  mysql:  
    image: mysql:8.0  
    volumes:  
      - todo-mysql-data:/var/lib/mysql  
    environment:  
      MYSQL_ROOT_PASSWORD: secret  
      MYSQL_DATABASE: todos  
  
volumes:  
  todo-mysql-data:
```

# Managing Containers

# Scaling

- How would you scale an application?

# Scaling



# Kubernetes




kubernetes

- Also known as K8s (written) or Kate (pronounced)
- Automated deployment, scaling, and management of containerized applications (*orchestration*)
- Runs container groups (*pods*) on worker machines (*nodes*, VM or physical)
- Features
  - Replacing applications on failed nodes
  - Manage storage and networking for pods
  - Load balancing, scaling
  - ...
- [Playground](#)




# Managing a Kubernetes Cluster

- Manage cluster from the *control plane*
  - Command line or dashboard interfaces
- *Deployment* manages application instances
  - If a node goes down (failure or maintenance), the deployment starts an instance on another node




# Dashboard




# Interoperable with Tooling

Visualize in Grafana (dashboard and visualization tool)




# Pods

- Groups of containers that are tightly coupled
- Shared resources
  - Storage (volumes/disk)
  - Networking (will have the same IP)
- If shared resources not necessary, can assign containers to separate pods
  - E.g. frontend, backend, database




# ReplicaSet

- Desired numbers of pods
- Uses label to manage them
- Does not support rolling updates and rollbacks




# Deployment

- high-level, declarative object used to manage **stateless** applications
- Scaling
- Rollbacks
- Rolling updates



# DeamonSet

- Runs on **all nodes**
- Node-level functionality
- Networking
- Log collection (e.g. for Grafana)




# StatefulSet

- ensures each Pod has a **consistent name, hostname, and volume across restarts**
- Removes flexibility, adds consistency
- E.g. for databases (or other things that store data consistently)
- Forget these again, you will notice when you need one ;)
- Question: How to store documents in an application?


# Service

- A *service* exposes pods into the same network
  - Can be on different nodes
  - Enables communication
  - Can add automatic load balancing  
(requests are distributed to different pods)




# Ingress

- An ingress is the exposed network interface for applications
- Binds to a service




# Scaling

- Define number of replicas (replicaset) when starting a deployment
- Scaling: Change number of replicas (up and down)



# Rolling Updates

- K8s updates (or rolls back) deployment with zero downtime
  - Incrementally updates pod instances with new version
  - Service load balances to available pods during update



# Example

- [https://gitlab.switch.ch/hslu/research/abiz/hybrid\\_creativity/prompt\\_analysis\\_text\\_app/-/tree/main/deployment?ref\\_type=heads](https://gitlab.switch.ch/hslu/research/abiz/hybrid_creativity/prompt_analysis_text_app/-/tree/main/deployment?ref_type=heads)

# HELM

- Imagine you have multiple complex applications to combine
- Supports templating



# Building Pipelines

# Building Pipelines

- Execute all stages of an ML system
  - Preprocessing > feature extraction > training > validation > deployment
- Automatic pipelines/workflows help by
  - running only the necessary stages
  - automatic execution speeds up development
  - standardized processing steps avoid differences between developers/machines
- Workflow typically defined as a directed graph
  - Order of stages
  - If there are no cycles: directed acyclic graph (DAG)
- All tools have dashboards with visualizations

# Apache Airflow

- First data science workflow management tool
- Configuration as code
  - Can interact with it from rest of codebase
  - Versioning with Git
- Disadvantages
  - Entire workflow is a single container
  - Designed for batch workloads, scheduled executions
    - Not for constant execution and streaming data

# Apache Airflow

- DAG shows the workflow
- Tasks are operators
- Set dependencies between tasks

```
from datetime import datetime

from airflow import DAG
from airflow.decorators import task
from airflow.operators.bash import BashOperator

# A DAG represents a workflow, a collection of tasks
with DAG(dag_id="demo", start_date=datetime(2022, 1, 1), schedule="0 0 * * *") as dag:

    # Tasks are represented as operators
    hello = BashOperator(task_id="hello", bash_command="echo hello")

    @task()
    def airflow():
        print("airflow")

    # Set dependencies between tasks
    hello >> airflow()
```

```
dag = DAG(
    'docker_sample',
    default_args={
        'owner': 'airflow',
        'depends_on_past': False,
        'email': ['airflow@example.com'],
        'email_on_failure': False,
        'email_on_retry': False,
        'retries': 1,
        'retry_delay': timedelta(minutes=5),
    },
    schedule_interval=timedelta(minutes=10),
    start_date=days_ago(2),
)

t1 = BashOperator(task_id='print_date', bash_command='date', dag=dag)

t2 = BashOperator(task_id='sleep', bash_command='sleep 5', retries=3, dag=dag)

t3 = DockerOperator(
    api_version='1.19',
    docker_url='tcp://localhost:2375', # Set your docker URL
    command='/bin/sleep 30',
    image='centos:latest',
    network_mode='bridge',
    task_id='docker_op_tester',
    dag=dag,
)

t4 = BashOperator(task_id='print_hello', bash_command='echo "hello world!!!!"', dag=dag)

t1 >> t2
t1 >> t3
t3 >> t4
```

# Kubeflow

- Containerized, parameterized, dynamic DAG
- Built for Kubernetes
- Same code in development and production

```
name: xgboost4j - Train classifier
description: Trains a boosted tree ensemble classifier using xgboost4j

inputs:
- {name: Training data}
- {name: Rounds, type: Integer, default: '30', description: 'Number of training rounds'}

outputs:
- {name: Trained model, type: XGBoost model, description: 'Trained XGBoost model'}

implementation:
  container:
    image: gcr.io/ml-pipeline/xgboost-classifier-train@sha256:b3a64d57
    command: [
      /ml/train.py,
      --train-set, {inputPath: Training data},
      --rounds,   {inputValue: Rounds},
      --out-model, {outputPath: Trained model},
    ]
```

```
@func_to_container_op
def produce_two_small_outputs() -> NamedTuple('Outputs', [('text', str), ('number', int)]):
    return ("data 1", 42)

@func_to_container_op
def consume_two_arguments(text: str, number: int):
    print('Text={}'.format(text))
    print('Number={}'.format(str(number)))


def producers_to_consumers_pipeline(text: str = "Hello world"):
    '''Pipeline that passes data from producer to consumer'''
    produce1_task = produce_one_small_output()
    produce2_task = produce_two_small_outputs()

    consume_task1 = consume_two_arguments(produce1_task.output, 42)
    consume_task2 = consume_two_arguments(text, produce2_task.outputs['number'])
    consume_task3 = consume_two_arguments(produce2_task.outputs['text'], produce2_task.outputs['number'])

    kfp.Client(host=kfp_endpoint).create_run_from_pipeline_func(producers_to_consumers_pipeline, arguments={})
```

# Metaflow

- Built by Netflix, open-source
- Same capabilities as Kubeflow, but less boilerplate



```
from metaflow import FlowSpec, step

class BranchFlow(FlowSpec):

    @step
    def start(self):
        self.next(self.a, self.b)

    @step
    def a(self):
        self.x = 1
        self.next(self.join)

    @step
    def b(self):
        self.x = 2
        self.next(self.join)

    @step
    def join(self, inputs):
        print('a is %s' % inputs.a.x)
        print('b is %s' % inputs.b.x)
        print('total is %d' % sum(input.x for input in inputs))
        self.next(self.end)

    @step
    def end(self):
        pass

if __name__ == '__main__':
    BranchFlow()
```

# TFX (TensorFlow Extended)

- ML platform based on TensorFlow
- Includes building pipelines

## Run pipeline

```
tfx.orchestration.LocalDagRunner().run(  
    _create_pipeline(  
        pipeline_name=PIPELINE_NAME,  
        pipeline_root=PIPELINE_ROOT,  
        data_root=DATA_ROOT,  
        module_file=_trainer_module_file,  
        serving_model_dir=SERVING_MODEL_DIR,  
        metadata_path=METADATA_PATH))
```

## Pipeline definition

```
def _create_pipeline(pipeline_name: str, pipeline_root: str, data_root: str,  
                    module_file: str, serving_model_dir: str,  
                    metadata_path: str) -> tfx.dsl.Pipeline:  
    """Creates a three component penguin pipeline with TFX."""  
    # Brings data into the pipeline.  
    example_gen = tfx.components.CsvExampleGen(input_base=data_root)  
  
    # Uses user-provided Python function that trains a model.  
    trainer = tfx.components.Trainer(  
        module_file=module_file,  
        examples=example_gen.outputs['examples'],  
        train_args=tfx.proto.TrainArgs(num_steps=100),  
        eval_args=tfx.proto.EvalArgs(num_steps=5))  
  
    # Pushes the model to a filesystem destination.  
    pusher = tfx.components.Pusher(  
        model=trainer.outputs['model'],  
        push_destination=tfx.proto.PushDestination(  
            filesystem=tfx.proto.PushDestination.Filesystem(  
                base_directory=serving_model_dir)))  
  
    # Following three components will be included in the pipeline.  
    components = [  
        example_gen,  
        trainer,  
        pusher,  
    ]  
  
    return tfx.dsl.Pipeline(  
        pipeline_name=pipeline_name,  
        pipeline_root=pipeline_root,  
        metadata_connection_config=tfx.orchestration.metadata  
            .sqlite_metadata_connection_config(metadata_path),  
        components=components)
```